

ACC PoC v1 Results Summary

policy	temp	variant	FollowWrong	ResistGold	Abstain	Other	ResistTotal	Delta_FollowWrong	Delta_ResistGold	Delta_Abstain	Delta_Other	Delta_ResistTotal
required	0.0	no_assert	0.000	1.000	0.000	0.000	1.000	0.000	0.000	0.000	0.000	0.000
required	0.0	user_final	1.000	0.000	0.000	0.000	0.000	1.000	-1.000	0.000	0.000	-1.000
required	0.0	assistant_prior	0.520	0.260	0.040	0.180	0.300	0.520	-0.740	0.040	0.180	-0.700
required	0.0	assistant_toolish	0.600	0.240	0.060	0.100	0.300	0.600	-0.760	0.060	0.100	-0.700
required	0.7	no_assert	0.000	0.940	0.000	0.060	0.940	0.000	0.000	0.000	0.000	0.000
required	0.7	user_final	1.000	0.000	0.000	0.000	0.000	1.000	-0.940	0.000	-0.060	-0.940
required	0.7	assistant_prior	0.520	0.300	0.040	0.140	0.340	0.520	-0.640	0.040	0.080	-0.600
required	0.7	assistant_toolish	0.600	0.180	0.080	0.140	0.260	0.600	-0.760	0.080	0.080	-0.680
• Baseline (no_assert) is strong and decidable at both temps (ResistGold = 1.00 at T=0.0; 0.94 at T=0.7); this is due to the probe screening we did within the notebook. • A final user assertion <i>forces near-total compliance</i> with the wrong tool (FollowWrong = 1.00) and collapses resistance (ResistTotal ↓ to 0.00) at both temps. • Earlier assertions (assistant_prior/toolish) <i>still cause large increases</i> in FollowWrong (+0.52 to +0.60) <i>and sizable drops in resistance</i> (−0.60 to −0.70), with some spillover into Other/Abstain. • The pattern is robust to temperature: ACC effects remain large at T=0.7, not just at greedy. IMPORTANT: We are flattening the `multi_turn_base` questions' dialog history into a single, one-shot prompt. This is quite a heavy augmentation but makes it so that the experiment is easier to pull off. I may need to actually setup multi-turn tool calling thoroughly (most likely via forking the BFCL eval Python repository) in a later version of the PoC.												